

Programmierung R - Exercise

Vectorization

May 10, SS 2022 || Hannah Behrens

Wir geben Impulse

Task 1

Your turn

1.1 Load the data set `Airquality3.csv` we have created at our second appointment from April 12, 2022 (Exercise *Data manipulation*).

1.2 Compute the mean of each numeric variable (use different `apply`- functions (`apply()`, `lapply()`, `sapply()` and `vapply()`)). How do the outputs differ?

1.3 Compute the median of Temperature dependent on the month.

Tasks 1.1 and 1.2 - answers

1.1

```
1 # read csv-file named "Airquality3.csv" (exercise from April 12, 2022)
2 airqu3 <- read.table("Airquality3.csv", sep = ",")
```

1.2

```
1 # the first six variables are numeric:
2 str(airqu3)
```

```
## 'data.frame':    112 obs. of  8 variables:
## $ Ozone      : int  41 36 12 18 23 19 8 16 11 14 ...
## $ Solar.R    : int  190 118 149 313 299 99 19 256 290 274 ...
## $ Wind       : num  7.4 8 12.6 11.5 8.6 13.8 20.1 9.7 9.2 10.9 ...
## $ Temperature: num  19.4 22.2 23.3 16.7 18.3 ...
## $ Month      : int  5 5 5 5 5 5 5 5 5 5 ...
## $ Day        : int  1 2 3 4 7 8 9 12 13 14 ...
## $ Month_name  : chr   "May" "May" "May" "May" ...
## $ Date       : chr   "1973-05-01" "1973-05-02" "1973-05-03" "1973-05-04" ...
```

Task 1.2 - apply() - answer

function apply()

```
1  # apply the mean function to the first six columns of airqu3
2  airqu3_apply <- apply(X = airqu3[,c(1:6)], MARGIN = 2, FUN = mean)
3  airqu3_apply
```

##	Ozone	Solar.R	Wind	Temperature	Month	Day
##	41.902	185.830	9.929	25.392	7.205	15.857

```
1  typeof(airqu3_apply) # output: a vector
```

```
## [1] "double"
```

Task 1.2 - lapply() - answer

```
1  airqu3_lapply <- lapply(X = airqu3[,c(1:6)], FUN = mean)
2  airqu3_lapply
```

```
## $Ozone
## [1] 41.9
##
## $Solar.R
## [1] 185.8
##
## $Wind
## [1] 9.929
##
## $Temperature
## [1] 25.39
##
## $Month
## [1] 7.205
##
## $Day
## [1] 15.86
```

function lapply(),
return value: a list

```
1  typeof(airqu3_lapply) # output: a list
```

```
## [1] "list"
```

Task 1.2 - answer

```
1 air_list <- list(Ozone = airqu3$Ozone, Solar.R = airqu3$Solar.R,  
2                 Wind = airqu3$Wind, Temperature = airqu3$Temperature,  
3                 Month = airqu3$Month, Day = airqu3$Day)  
4 air_list
```

Saving airqu3 as a list named air_list

```
## $Ozone  
## [1] 41 36 12 18 23 19 8 16 11 14 18 14 34 6 30 11 1 11  
## [19] 4 32 23 45 115 37 29 71 39 23 21 37 20 12 13 135 49 32  
## [37] 64 40 77 97 97 85 10 27 7 48 35 61 79 63 16 80 108 20  
## [55] 52 82 50 64 59 39 9 16 122 89 110 44 28 65 22 59 23 31  
## [73] 44 21 9 45 168 73 76 118 84 85 96 78 73 91 47 32 20 23  
## [91] 21 24 44 21 28 9 13 46 18 13 24 16 13 23 36 7 14 30  
## [109] 14 18 20 20  
##  
## $Solar.R  
## [1] 190 118 149 313 299 99 19 256 290 274 65 334 307 78 322 44 8 320  
## [19] 25 92 13 252 223 279 127 291 323 148 191 284 37 120 137 269 248 236  
## [37] 175 314 276 267 272 175 264 175 48 260 274 285 187 220 7 294 223 81  
## [55] 82 213 275 253 254 83 24 77 255 229 207 192 273 157 71 51 115 244  
## [73] 190 259 36 212 238 215 203 225 237 188 167 197 183 189 95 92 252 220  
## [91] 230 259 236 259 238 24 112 237 224 27 238 201 238 14 139 49 20 193  
## [109] 191 131 223 300  
##  
## $Wind  
## [1] 7.4 8.0 12.6 11.5 8.6 13.8 20.1 9.7 9.2 10.9 13.2 11.5 12.0 18.4 11.5  
## [16] 9.7 9.7 16.6 9.7 12.0 12.0 14.9 5.7 7.4 9.7 13.8 11.5 8.0 14.9 20.7
```

Task 1.2 - lapply() - answer

```
1  airqu3_lapply2 <- lapply(X = air_list, FUN = mean)
2  airqu3_lapply2
```

```
## $Ozone
## [1] 41.9
##
## $Solar.R
## [1] 185.8
##
## $Wind
## [1] 9.929
##
## $Temperature
## [1] 25.39
##
## $Month
## [1] 7.205
##
## $Day
## [1] 15.86
```

applying lapply() to a list → a list is returned

```
1  typeof(airqu3_lapply2) # output: a list
```

```
## [1] "list"
```

Task 1.2 - sapply() - answer

```
function sapply()
```

```
1  airqu3_sapply <- sapply(X = airqu3[,c(1:6)], FUN = mean)
2  airqu3_sapply
```

##	Ozone	Solar.R	Wind	Temperature	Month	Day
##	41.902	185.830	9.929	25.392	7.205	15.857

```
1  typeof(airqu3_sapply) # output: a vector
```

```
## [1] "double"
```


Task 1.2 - sapply() - answer

applying sapply() to a list → a vector is returned when simplify = TRUE

```
1 airqu3_sapply2 <- sapply(air_list, FUN = mean)
2 airqu3_sapply2
```

##	Ozone	Solar.R	Wind	Temperature	Month	Day
##	41.902	185.830	9.929	25.392	7.205	15.857

```
1 typeof(airqu3_sapply2) # output: a vector
```

```
## [1] "double"
```

What happens when simplify is set to FALSE?

Task 1.2 - sapply() - answer

```
1 airqu3_sapply3 <- sapply(air_list, FUN = mean, simplify = FALSE)
2 airqu3_sapply3
```

```
## $Ozone
## [1] 41.9
##
## $Solar.R
## [1] 185.8
##
## $Wind
## [1] 9.929
##
## $Temperature
## [1] 25.39
##
## $Month
## [1] 7.205
##
## $Day
## [1] 15.86
```

Setting simplify = FALSE: a list is returned

```
1 typeof(airqu3_sapply3) # output: a vector
```

```
## [1] "list"
```

Task 1.2 - vapply() - answer

function vapply(): specification of the return value (FUN.VALUE) is necessary

```
1 airqu3_vapply <- vapply(X = airqu3[,c(1:6)], FUN = mean, FUN.VALUE = numeric(1))
2 airqu3_vapply
```

##	Ozone	Solar.R	Wind	Temperature	Month	Day
##	41.902	185.830	9.929	25.392	7.205	15.857

```
1 typeof(airqu3_vapply) # output: a vector
```

```
## [1] "double"
```

Task 1.3 - answer

1.3 Based on Airquality3.csv, compute the median of Temperature dependent on the month.

```
1 tapply(X = airqu3$Temperature, INDEX = airqu3$Month, FUN = median)
```

```
##      5      6      7      8      9  
## 18.89 24.72 28.89 27.22 24.44
```

```
1 # or
```

```
2 tapply(X = airqu3$Temperature, INDEX = airqu3$Month_name, FUN = median)
```

```
##      August      July      June      May September  
##      27.22      28.89      24.72      18.89      24.44
```

```
1 # --> is a big advantage compared to calculating the median of the temperature  
2 # for each month piecewise like  
3 median(airqu3$Temperature[which(airqu3$Month == 5)])
```

```
## [1] 18.89
```

```
1 # ... and so on
```

Task 2

Remember the data set `datasets::airquality` (R Core Team 2021a). It consists of the following variables

```
1 colnames(airquality)
```

```
## [1] "Ozone" "Solar.R" "Wind" "Temp" "Month" "Day"
```

which are numeric.

We are interested in all possible and different linear models when one of these six variables is considered as the dependent variable y . How can we achieve this?

How many combinations of different linear models exist for each y ?

Task 2 - number of combinations - answer

For each y:

- $N = 5$ remaining variables as predictors to form linear models, e.g. $y = \text{Ozone}$:

```
1 lm(Ozone ~ 1 + Solar.R + Wind + Temp + Month + Day)
2 # but also
3 lm(Ozone ~ 1 + Solar.R + Wind + Temp + Month)
4 # as well as
5 lm(Ozone ~ 1 + Solar.R + Wind + Temp + Day)
6 # and so on...
```

- As we have seen, a linear model for Ozone can be created based on 5, 4, 3, 2 and 1 variable(s) as predictor(s).
- Each predictor appears only **once** in each linear model and the ordering **does not** play any role: So, the number of combinations when picking up 5 of 5, 4 of 5 predictors and so on is calculated based on the binomial coefficient $\binom{N}{n} = \frac{N!}{n! \cdot (N-n)!}$. (Fahrmeir et al. 2016, 187f.)

Task 2 - number of combinations - answer

From that it follows:

- When forming a model with 5 predictors, only 1 combination is possible:

$$\binom{5}{5} = \frac{5!}{5! \cdot (5-5)!} = 1$$

- When forming a model with 4 predictors, 5 combinations are possible:

$$\binom{5}{4} = \frac{5!}{4! \cdot (5-4)!} = 5$$

→ So, for each y

```
1 combi_per_y <- choose(5,5) + choose(5,4) + choose(5,3) + choose(5,2) + choose(5,1)
2 combi_per_y
```

```
## [1] 31
```

different linear models can be created and examined.

Task 2 - implementation - step 1

Your turn

Based on the presented theory, the first step is to *combine the predictors* to create 31 linear models in the next step. These *combinations* should be saved in a list. How can we do this?

Task 2 - implementation - step 1 - answer

Combining the predictors to create 31 linear models

```
1  ## main program
2  data <- airquality
3  coln_names <- colnames(data) # get column names
4
5  y <- coln_names[1] # e.g. Ozone (column 1) as the dependent variable (y)
6  vars_grid <- coln_names[-1] # the remaining variables are predictors
7  variables_grid <- expand.grid(replicate(length(vars_grid), vars_grid,
8                                simplify = FALSE)) # make a grid with all possible combinations
9  # since each predictor can only be listed once in the linear model, remove replicates
10 variables_grid_unique <- apply(variables_grid, MARGIN = 1, FUN = unique)
11 # sort entries of the sublists since the order of the predictors does not play any role
12 variables_grid_sort <- lapply(variables_grid_unique, sort)
13 # then remove replicates again
14 variables_grid_sort_unique <- unique(variables_grid_sort)
15 # add y to the end of all sublists:
16 variables_grid_sort_unique_y <- lapply(variables_grid_sort_unique,
17                                         FUN = function(x){c(x, x[length(x)+1] <- y)})
18 air_grid_lm <- variables_grid_sort_unique_y
19 # ... something is missing here
```

What to do next? Any ideas? Suggestions?

Your turn

Based on our list `air_grid_lm`, we have to create the formulas of each linear model with `y` (here: `Ozone`) as dependent variable. How can you implement it? Do it!

Task 2 - step 2 - implementation - answer

```
1  # x: the first element(s) of each sublist (= predictors)
2  # y: the last element of each sublist
3  determiningLinMod <- function(x, y = x[length(x)], dataset = airquality){
4    # predictors x1 ,..., x_{length(x)-1}
5    predictors <- paste(x[-length(x)], collapse = " + ") # concatenate the predictors
6    # by a + -sign
7    make_formula <- paste(y, " ~ 1 + ", predictors)
8    formula_lm <- as.formula(make_formula) # convert the string to a formula
9    lin_model <- lm(formula_lm, data = dataset) # call lm()
10   # return the different formulas as well as the model
11   return(list(formula = formula_lm, model = lin_model))
12 }
```

```
1  # apply the function determiningLinMod() to each combination listed in air_grid_lm:
2  lapply_lm_air <- lapply(X = air_grid_lm, FUN = determiningLinMod)
```

Task 2 - step 2 - implementation - answer

```
1  ## main program
2  data <- airquality
3  coln_names <- colnames(data) # get column names
4
5  y <- coln_names[1] # e.g. Ozone (column 1) as the dependent variable (y)
6  vars_grid <- coln_names[-1] # the remaining variables are predictors
7  variables_grid <- expand.grid(replicate(length(vars_grid), vars_grid,
8                                simplify = FALSE)) # make a grid with all possible combinations
9    # since each predictor can only be listed once in the linear model, remove replicates
10 variables_grid_unique <- apply(variables_grid, MARGIN = 1, FUN = unique)
11 # sort entries of the sublists since the order of the predictors does not play any role
12 variables_grid_sort <- lapply(variables_grid_unique, sort)
13    # then remove replicates again
14 variables_grid_sort_unique <- unique(variables_grid_sort)
15    # add y to the end of all sublists:
16 variables_grid_sort_unique_y <- lapply(variables_grid_sort_unique,
17                                         FUN = function(x){c(x, x[length(x)+1] <- y)})
18 air_grid_lm <- variables_grid_sort_unique_y
19 # insert:
20 lapply_lm_air <- lapply(air_grid_lm, FUN = determiningLinMod)
```

Task 2 - step 2 - implementation - results

```
1 length(lapply_lm_air) # 31 combinations expected (see some slides before)
```

```
## [1] 31
```

```
1 lapply_lm_air
```

```
## [[1]]
```

```
## [[1]]$formula
```

```
## Ozone ~ 1 + Solar.R
```

```
## <environment: 0x00000000255a5cd0>
```

```
##
```

```
## [[1]]$model
```

```
##
```

```
## Call:
```

```
## lm(formula = formula_lm, data = data)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      Solar.R
```

```
##      18.599         0.127
```

```
##
```

```
##
```

```
##
```

```
## [[2]]
```

```
## [[2]]$formula
```

```
## Ozone ~ 1 + Solar.R + Wind
```

Task 2 - step 3 - parallelization

Very nice! We get all possible linear models for the variable `Ozone` of the `airquality` data set (R Core Team 2021a).

But now, we want to go a step further by parallelizing the determination of the 31 linear models for `Ozone`.

How can we achieve this?

Task 2 - step 3 - parallelization

The function `parallel::mclapply()` (R Core Team 2021b) is available which works as follows (The example is taken from Peng (2020)):

```
1  r <- mclapply(1:10, function(i) {  
2      Sys.sleep(10)  # Do nothing for 10 seconds  
3  }, mc.cores = 1)  
4  r
```

```
## [[1]]  
## NULL  
##  
## [[2]]  
## NULL  
##  
## [[3]]  
## NULL  
##  
## [[4]]  
## NULL  
##  
## [[5]]  
## NULL  
##  
## [[6]]  
## NULL
```

Task 2 - step 3 - parallelization

```
1 parallelized_lm_air <- parallel::mclapply(X = air_grid_lm,  
2                                           FUN = determiningLinMod, mc.cores = 1)  
3 parallelized_lm_air
```

```
## [[1]]  
## [[1]]$formula  
## Ozone ~ 1 + Solar.R  
## <environment: 0x00000000279626e0>  
##  
## [[1]]$model  
##  
## Call:  
## lm(formula = formula_lm, data = data)  
##  
## Coefficients:  
## (Intercept)      Solar.R  
##      18.599        0.127  
##  
##  
##  
## [[2]]  
## [[2]]$formula  
## Ozone ~ 1 + Solar.R + Wind  
## <environment: 0x000000002799f010>  
##  
## [[2]]$model
```

Instead of using `lapply()` `mclapply` (R Core Team 2021b) is used.
Parallelization based on `mclapply` is not possible on Windows.

Task 2 - step 3 - parallelization

```
1 length(parallelized_lm_air)
```

```
## [1] 31
```

The expectations have been fulfilled!

Task 3

The following list is given:

```
1 L <- list(money = c(250, 124, 360, 720, 340, 340),  
2           hours = c(19, 12, 30, 48, 26, 25),  
3           idx = c(1:6),  
4           name = c("Paul", "Emma", "Mia", "John", "Kim", "Maxi"))
```

Your turn

Make use of the `*apply()`-functions and

- a) compute the total sum of `money`, `hours` and `idx` and return a list.
- b) compute the mean of `money` and `hours` and return a vector.
- c) multiply each element of `money` and `hours` by 1.5 and return a matrix.

Tasks 3a) and 3b) - answers

```
1 # 3.a)
2 lapply(X = L[1:3], FUN = sum)
```

```
## $money
## [1] 2134
##
## $hours
## [1] 160
##
## $idx
## [1] 21
```

```
1 # 3.b)
2 sapply(X = L[1:2], FUN = mean)
```

```
## money hours
## 355.67 26.67
```

```
1 # vapply(L[1:2], mean, FUN.VALUE=numeric(1)) # alternative
2 # --> the value specified by FUN.VALUE (here: numeric(1)) and the actual value returned
3 # (here: the mean of money and hours) have to match
```

Task 3.c) - answer

```
1  sapply(X = L[1:2], FUN = function(x){x * 1.5})
```

```
##      money hours
## [1,]   375  28.5
## [2,]   186  18.0
## [3,]   540  45.0
## [4,]  1080  72.0
## [5,]   510  39.0
## [6,]   510  37.5
```

Document the code (main program (slide 23) and function `determiningLinMod()`) as you have learned in the last lecture and exercise about *Software Development*.

You can also think of wrapping the code from slide 23 in a function.

- Buchwitz, B. 2021. *Computational Statistics*.
<https://bchwtz.github.io/bchwtz-cswr/>.
- Fahrmeir, L., C. Heumann, R. Künstler, I. Pigeot, and G. Tutz. 2016. *Statistik. Der Weg Zur Datenanalyse*. 8th ed. Springer-Lehrbuch. Berlin: Springer.
<https://doi.org/10.1007/978-3-662-50372-0>.
- Peng, R. D. 2020. *R Programming for Data Science*. 22.3 the Parallel Package. Vienna, Austria: R Foundation for Statistical Computing.
<https://bookdown.org/rdpeng/rprogdatascience/parallel-computation.html#the-parallel-package>.
- R Core Team. 2021a. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing.
<https://www.R-project.org/>.
- . 2021b. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing.
<https://www.R-project.org/>.